LATENTSIG REVIEW ASSIGNMENT · DEMAND FORECASTING AGENT WORKFLOW

Demand Forecasting Agent

LightGBM demand forecasts on Walmart M5 data, forecast-driven stockout detection, supplier allocation and human-approved purchase orders, orchestrated by an agent.

Kaggle M5 | LightGBM | Agentic workflow · 9 tools | FastAPI + Next.js | 48 tests

1. The business problem

A retail supply chain re-orders stock on a fixed weekly cycle using a static rule: each week, order enough to cover the next review period and the supplier lead time at the recent average sales rate. This is the "traditional static system" the brief describes. It has three blind spots:

1. **It looks backwards.** A trailing average does not see SNAP benefit days, holidays or promotions that are already on the calendar.
2. **It has no safety buffer.** Normal demand variability and short supplier shipments cause stockouts. Replaying this policy over the real M5 sales history runs a series out of stock on **2.4% of days**.
3. **It needs a planner to catch problems.** Someone has to check 15 store/department series every day, decide which ones need an extra order and split that order across suppliers.

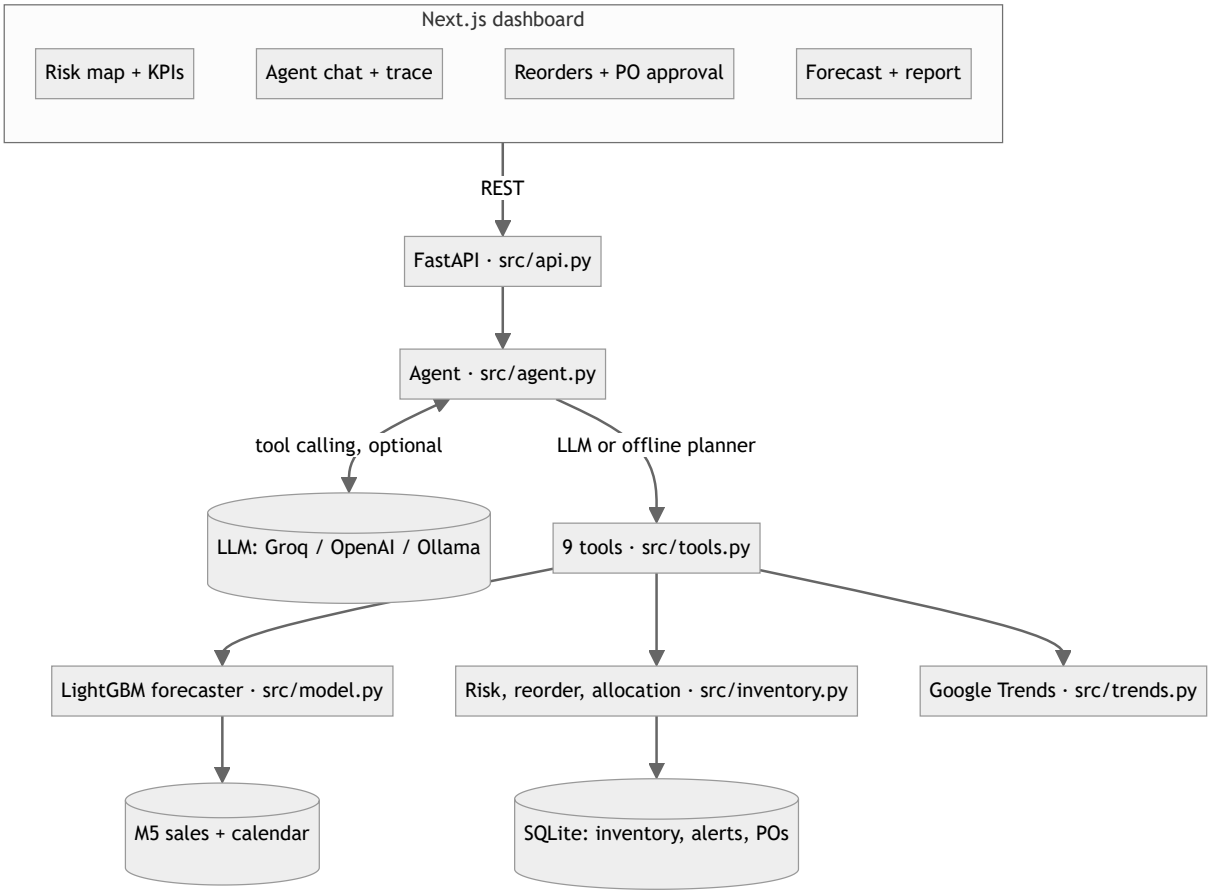
The agent automates that daily check. It forecasts demand with an ML model, finds the series that will run out before the next scheduled delivery, explains why (promotions, spikes, external trends), sizes an expedite order, splits it across suppliers, and drafts purchase orders for a human to approve.

2. Requirement coverage

Brief asks for	Where it is
ML tool: LightGBM or Prophet on M5 / Walmart	`src/model.py` : global **LightGBM** model trained on **Kaggle M5** sales
Detect sudden stockout risks	`detect_stockout_risk` , `run_risk_scan` (`src/inventory.py`)
Detect demand spikes	`detect_demand_spike` : last 7 days *and* next 7 forecast days vs a 90-day baseline
Query external trend data	`query_trends` : live Google Trends via `pytrends` , cached, honest "unavailable" fallback
Query promotional schedules	`get_promo_calendar` : SNAP days and retail events from the M5 calendar (56 days ahead)
Recommend automated re-orders	`recommend_reorder` and `draft_purchase_order` , with human approval in the UI
Supplier allocations	`allocate_suppliers` : urgent units go to the fastest supplier, the rest to the best cost/reliability supplier that can still deliver in time

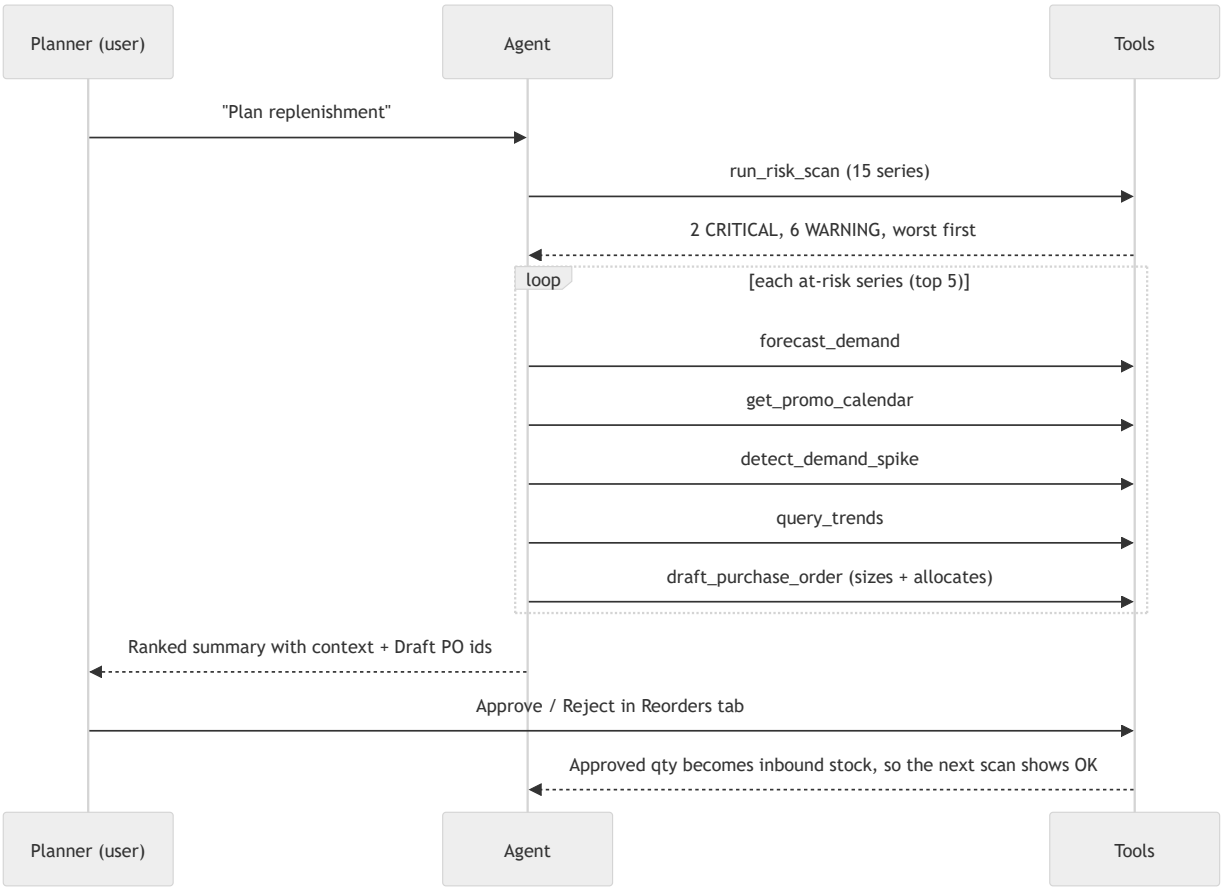
Brief asks for	Where it is
Report on architecture and workflow	this document and `docs/ARCHITECTURE.md`
Bonus: UI/UX to run the agent	Next.js dashboard (`frontend/`): risk map, agent chat with reasoning trace, forecast chart, reorders and PO approval, report
AI tools used	Section 10

3. Architecture



Planners. With an API key (Groq is free) the LLM chooses the tools using OpenAI-compatible function calling, for up to 8 tool steps. Without a key, an offline planner detects the intent and runs fixed multi-step plans with the same tools. Either way, every call appears in the UI's reasoning trace with its thought, arguments and full observation.

4. The agentic workflow ("Plan replenishment")



Example output (offline planner, real data):

> **CA_1/HOBBIES [CRITICAL]** position 6640 vs 6817 forecast demand to next delivery, projected stockout 2016-05-07. Order 717 units: 717 from CraftHub (10d). Draft PO #1 awaiting approval. Context: 8 SNAP days in next 2 weeks; events: Pesach End, OrthodoxEaster, Cinco De Mayo; Google Trends 'toys' down (-14.4% vs 12-month average).

The loop closes: approving the drafted orders adds them to inbound stock. In testing, approving the five drafts moved the portfolio from 8 at-risk series (2 critical) to 3 (0 critical), and the next run of the agent picked up the remaining three.

5. Data

- **Source:** Kaggle *M5 Forecasting – Accuracy* (`sales_train_validation.csv`, `calendar.csv`). Item-level sales are aggregated to **5 stores × 3 departments = 15 daily series**, 2011-01-29 to 2016-04-24 (28,695 rows).
- **Calendar features:** SNAP days (per state) as the promotion flag, and M5 events as the holiday flag. The calendar runs 56 days past the last sales day, so these inputs are genuinely known for the whole forecast horizon. They are saved to `data/m5_calendar_future.csv` so a fresh clone works without the raw M5 download.

- **Not in M5 (documented assumptions):**
  - *Inventory.* We replay the static weekly policy from Section 1 over actual sales. Each store reviews on its own weekday, and orders arrive after the supplier lead time. Shipments are short (40–90%) with probability 1 – supplier reliability. This produces on-hand stock, inbound stock and days-to-next-review for every series.
  - *Suppliers.* Two per department, each with a lead time, relative unit cost and on-time reliability (`src/data.py`). Faster suppliers are more expensive.

6. Forecasting model (the ML tool)

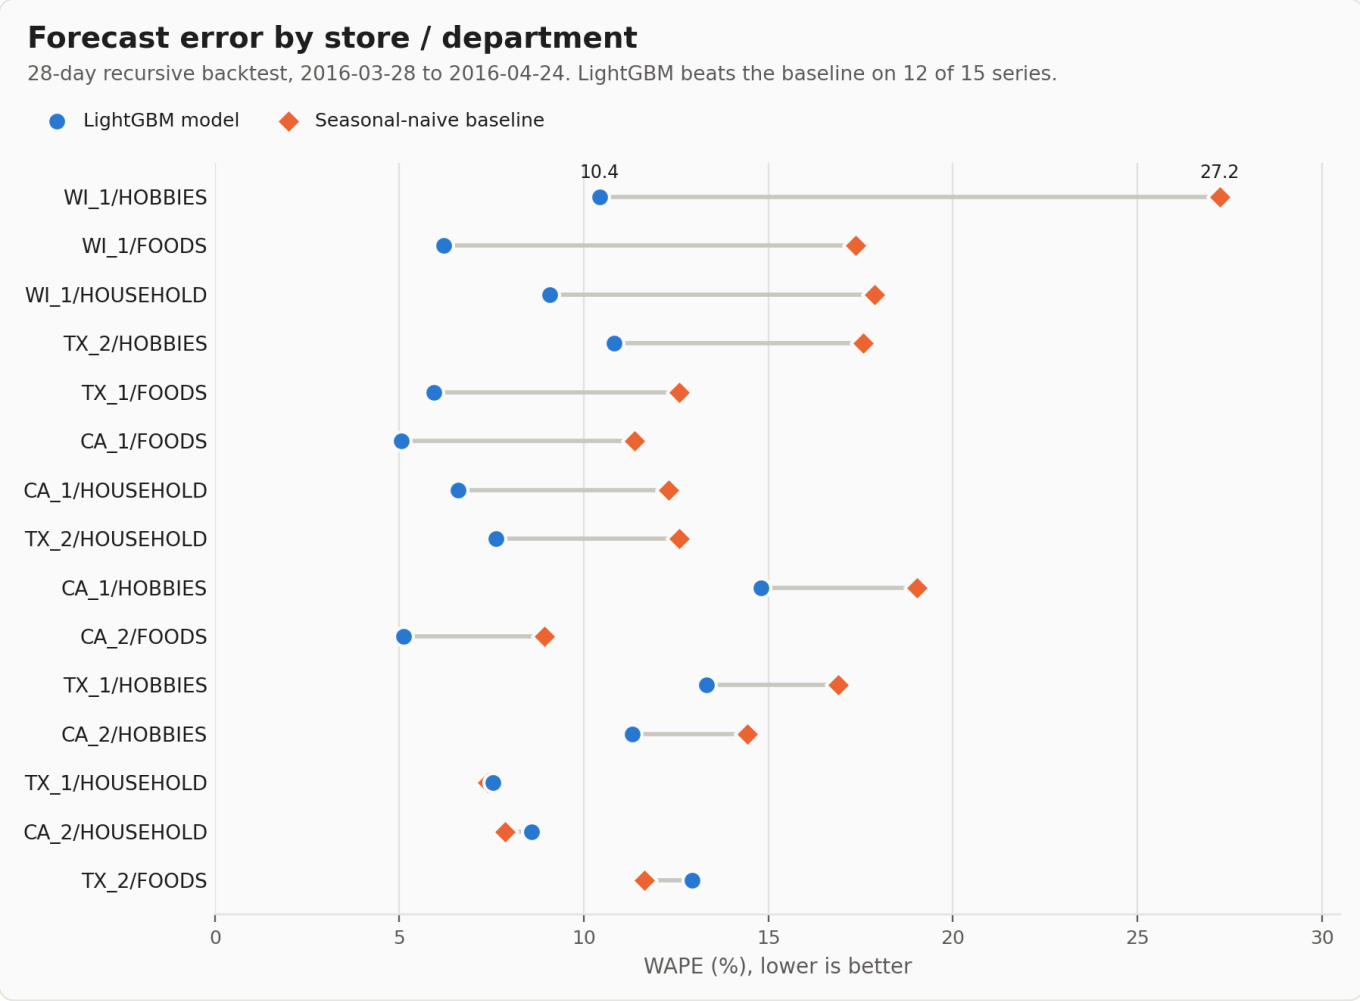
- **LightGBM** `LGBMRegressor`: one global model over all 15 series, 400 trees, learning rate 0.05, 31 leaves.
- **Features:** store, department, day of week, month, SNAP promotion, event/holiday, lags of 1, 7 and 28 days, and 7- and 28-day rolling means.
- **Recursive multi-step forecasting:** each predicted day feeds the next day's lags. Training and serving use the same function (`recursive_forecast`), so the backtest measures exactly what the app serves.
- **Prediction intervals:** an 80% empirical band per series, from the 10th and 90th percentiles of actual/forecast ratios in the backtest. These are real multi-step errors, not a fixed ± percentage.

Evaluation: chronological backtest, 28-day horizon, 3 rolling origins

Every window is trained only on data before its test start, and future sales are never used as inputs (`src/eval.py`, `reports/backtest_metrics.json`). Metrics are macro averages over the 15 series.

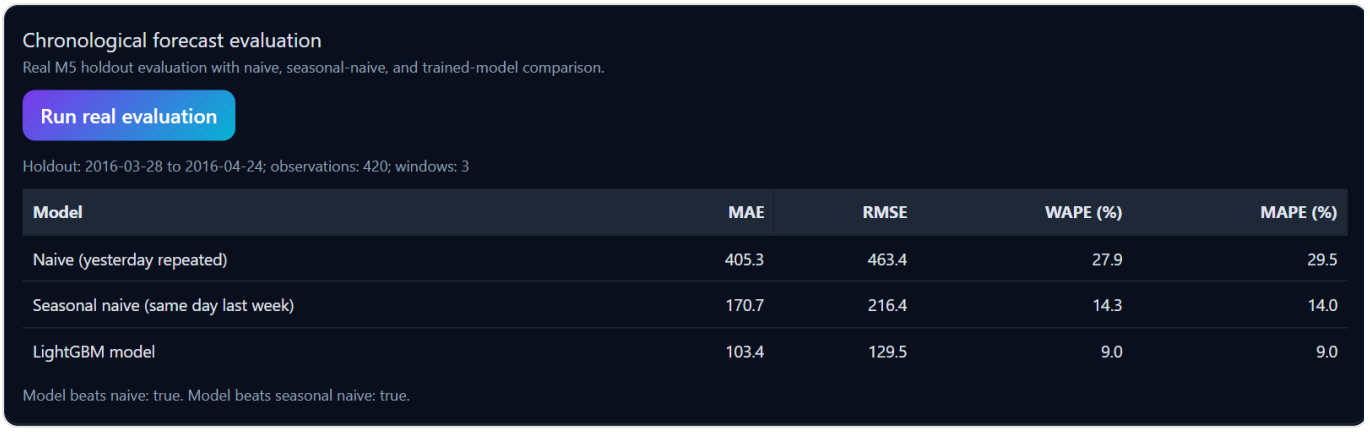
Window (test period)	Naive MAE / WAPE	Seasonal-naive MAE / WAPE	LightGBM MAE / WAPE
2016-03-28 → 04-24	405.3 / 27.9%	170.7 / 14.3%	**103.4 / 9.0%**
2016-02-29 → 03-27	461.2 / 36.5%	150.8 / 12.8%	**129.6 / 10.8%**
2016-02-01 → 02-28	424.4 / 32.9%	168.5 / 13.0%	**95.6 / 8.6%**
Mean	430.3 / 32.4%	163.4 / 13.4%	**109.5 / 9.5%**

LightGBM cuts error by about 30% compared with the seasonal-naive baseline, and wins in all three windows. Per series it wins on 12 of 15. It loses narrowly on CA_2/HOUSEHOLD and TX_1/HOUSEHOLD, and on TX_2/FOODS (12.9% vs 11.6% WAPE), which is the first thing to investigate next. HOBBIES is the hardest department (10–15% WAPE): it has low volume and higher noise.



Most recent 28-day window. Each row is one series; the grey bar is the error gap between the baseline (orange diamond) and LightGBM (blue circle). The bottom three rows are the series where the baseline is slightly better. Regenerate with `python scripts/plot_accuracy.py`.

The same results as they appear in the app's Report tab:



7. Risk, reorder and allocation logic

Because the store orders weekly, stock has to last until the order placed at the next review arrives. This window is the *protection interval*.

```
inventory_position = on_hand + on_order
protection_days    = days_to_next_review + supplier_lead_time
protection_demand  = Σ LightGBM forecast over protection_days
safety_stock       = z · σ_daily · √protection_days
reorder_point      = protection_demand + safety_stock

CRITICAL  position < protection_demand    → runs out before the next scheduled delivery
WARNING   position < reorder_point        → will dip into safety stock
OK        otherwise
```

Here z = 1.65 targets about a 95% cycle service level, and σ_daily is estimated from week-over-week changes in sales, so the regular weekly pattern doesn't inflate it.

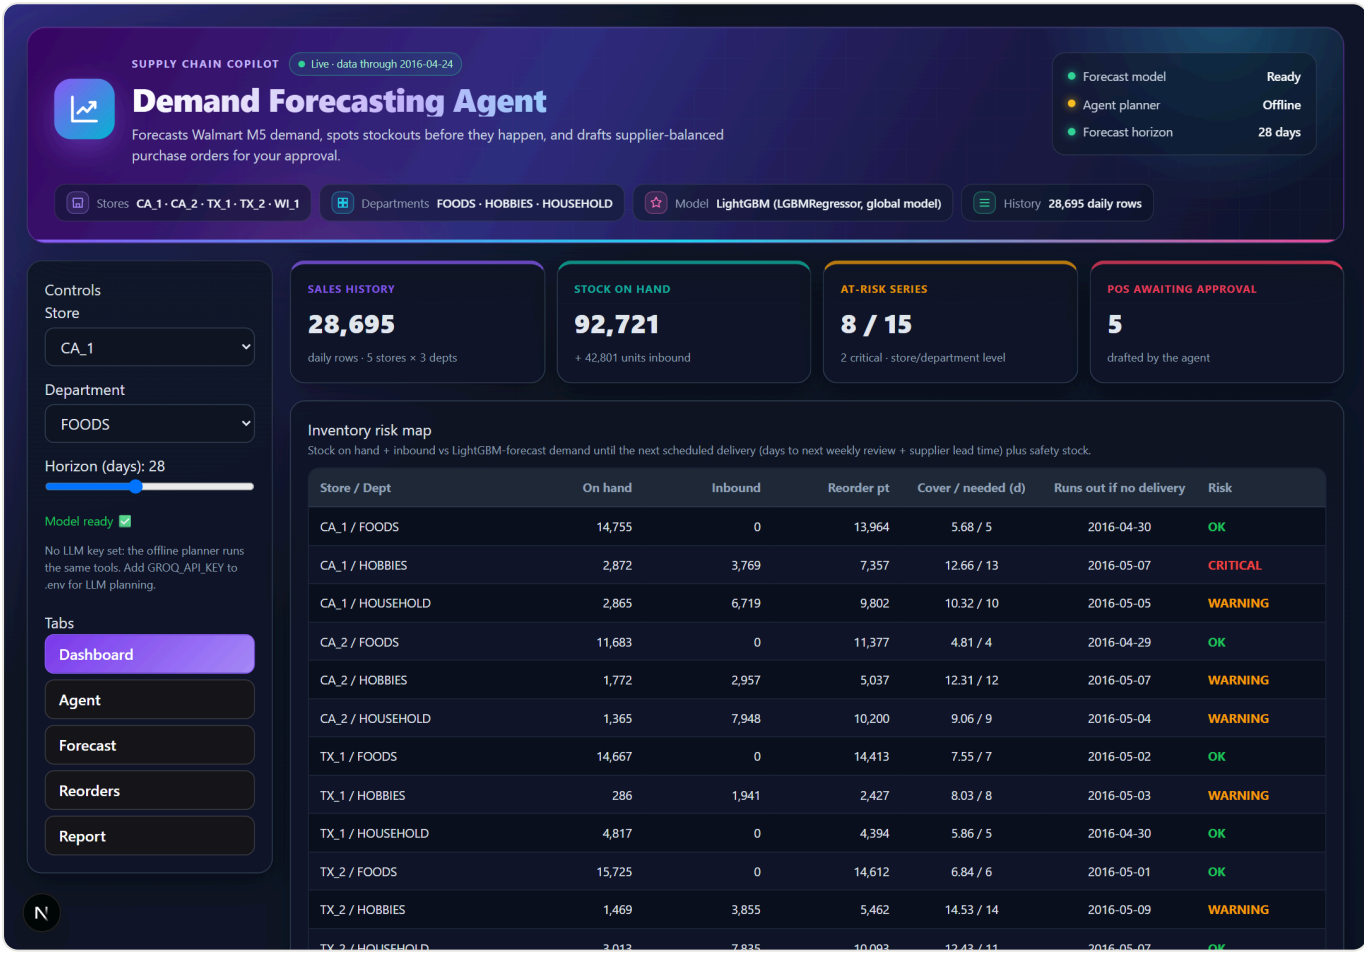
- **Expedite quantity:** `reorder_point - position`, rounded up. This tops the series back up so the routine weekly order takes over again. OK series get no order.
- **Supplier split:** units needed before even the fastest supplier can deliver go to the fastest supplier. The remainder goes to the supplier with the best unit cost ÷ reliability score, provided its lead time still fits the protection interval.
- **Purchase orders** are created as `DRAFT` with an idempotency key, so re-running the agent never duplicates them. Only a person can mark them `APPROVED` or `REJECTED`, and approved quantities become inbound stock.

8. UI/UX

A Next.js 16 + FastAPI app. All screenshots are from the running app on real M5 data.

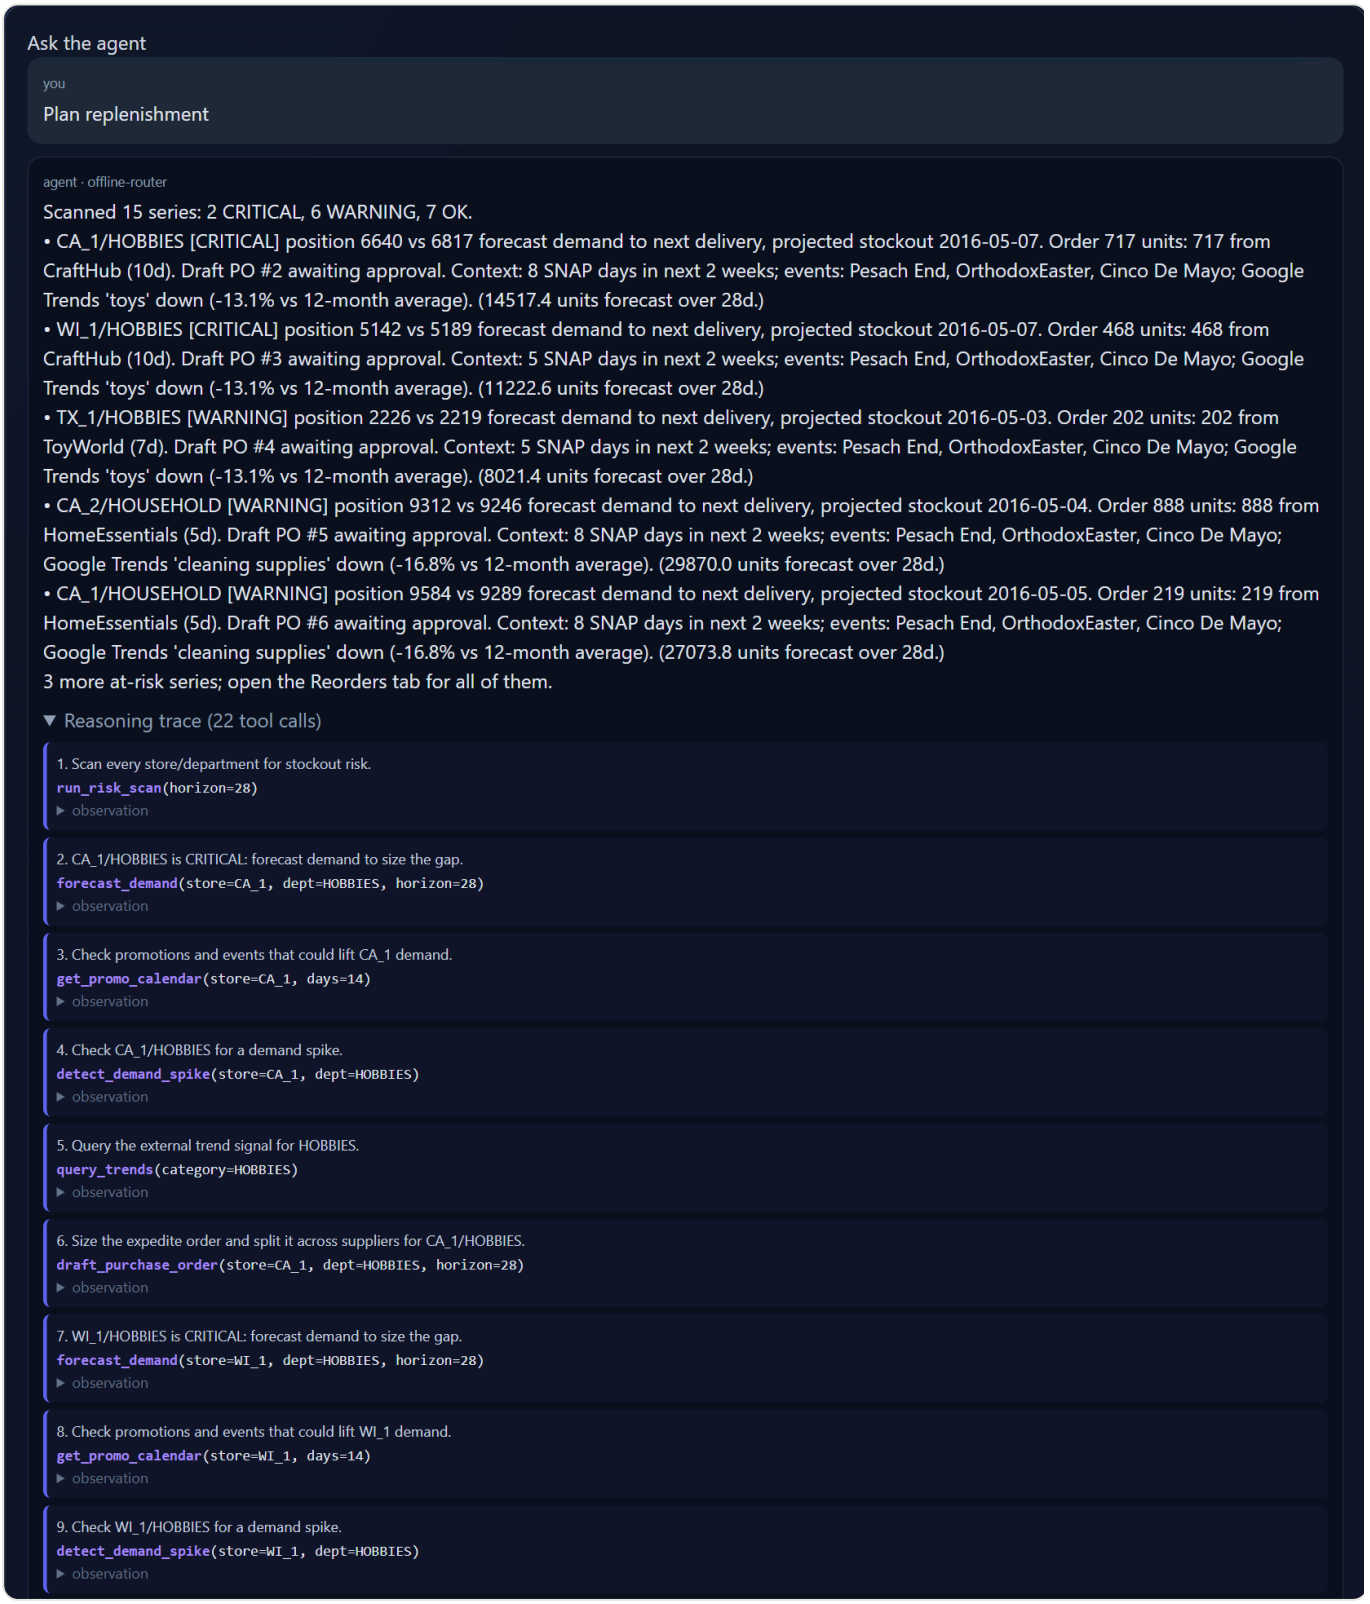
Dashboard

KPI cards, model and planner status, and a risk map for all 15 series: on hand, inbound, reorder point, days of cover vs days needed, and the date stock runs out with no delivery. Below it sit a risk feed with a one-click "Ask the agent to plan replenishment" button, a live Google Trends widget, and a recent-demand chart.



Agent

Chat with suggestion chips. Every answer has an expandable reasoning trace showing each step's thought, the tool called with its arguments, and the full observation. Here, "Plan replenishment" runs 22 tool calls across the at-risk series.



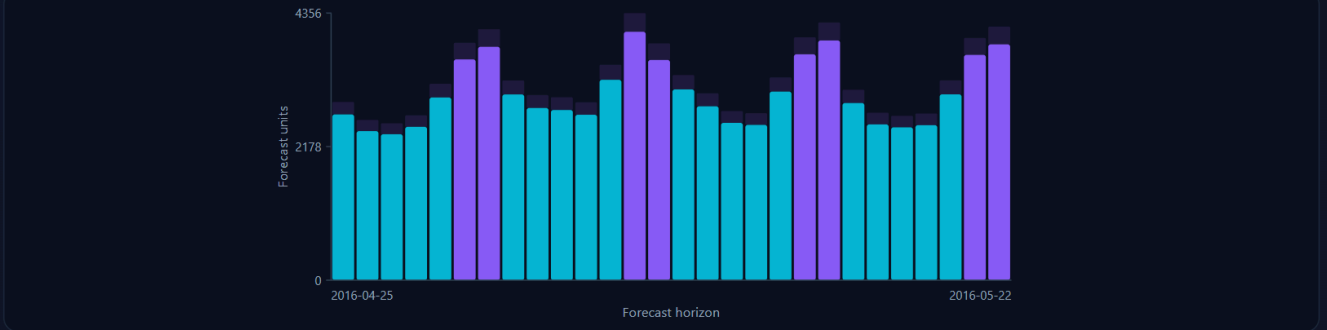
Forecast

A 7–56-day forecast. The darker cap on each bar is the empirical 80% upper band; the backtest error for this series appears next to the seasonal-naive baseline. Moving the horizon slider shows an "Updating forecast" overlay until the new forecast arrives.

Forecast · CA_1 / FOODS · 28 days

Total: **85327.6** Avg daily: **3047.4** Range: 2016-04-25 → 2016-05-22 Backtest MAPE 5.4% · MAE 153 (seasonal-naive MAE 344)

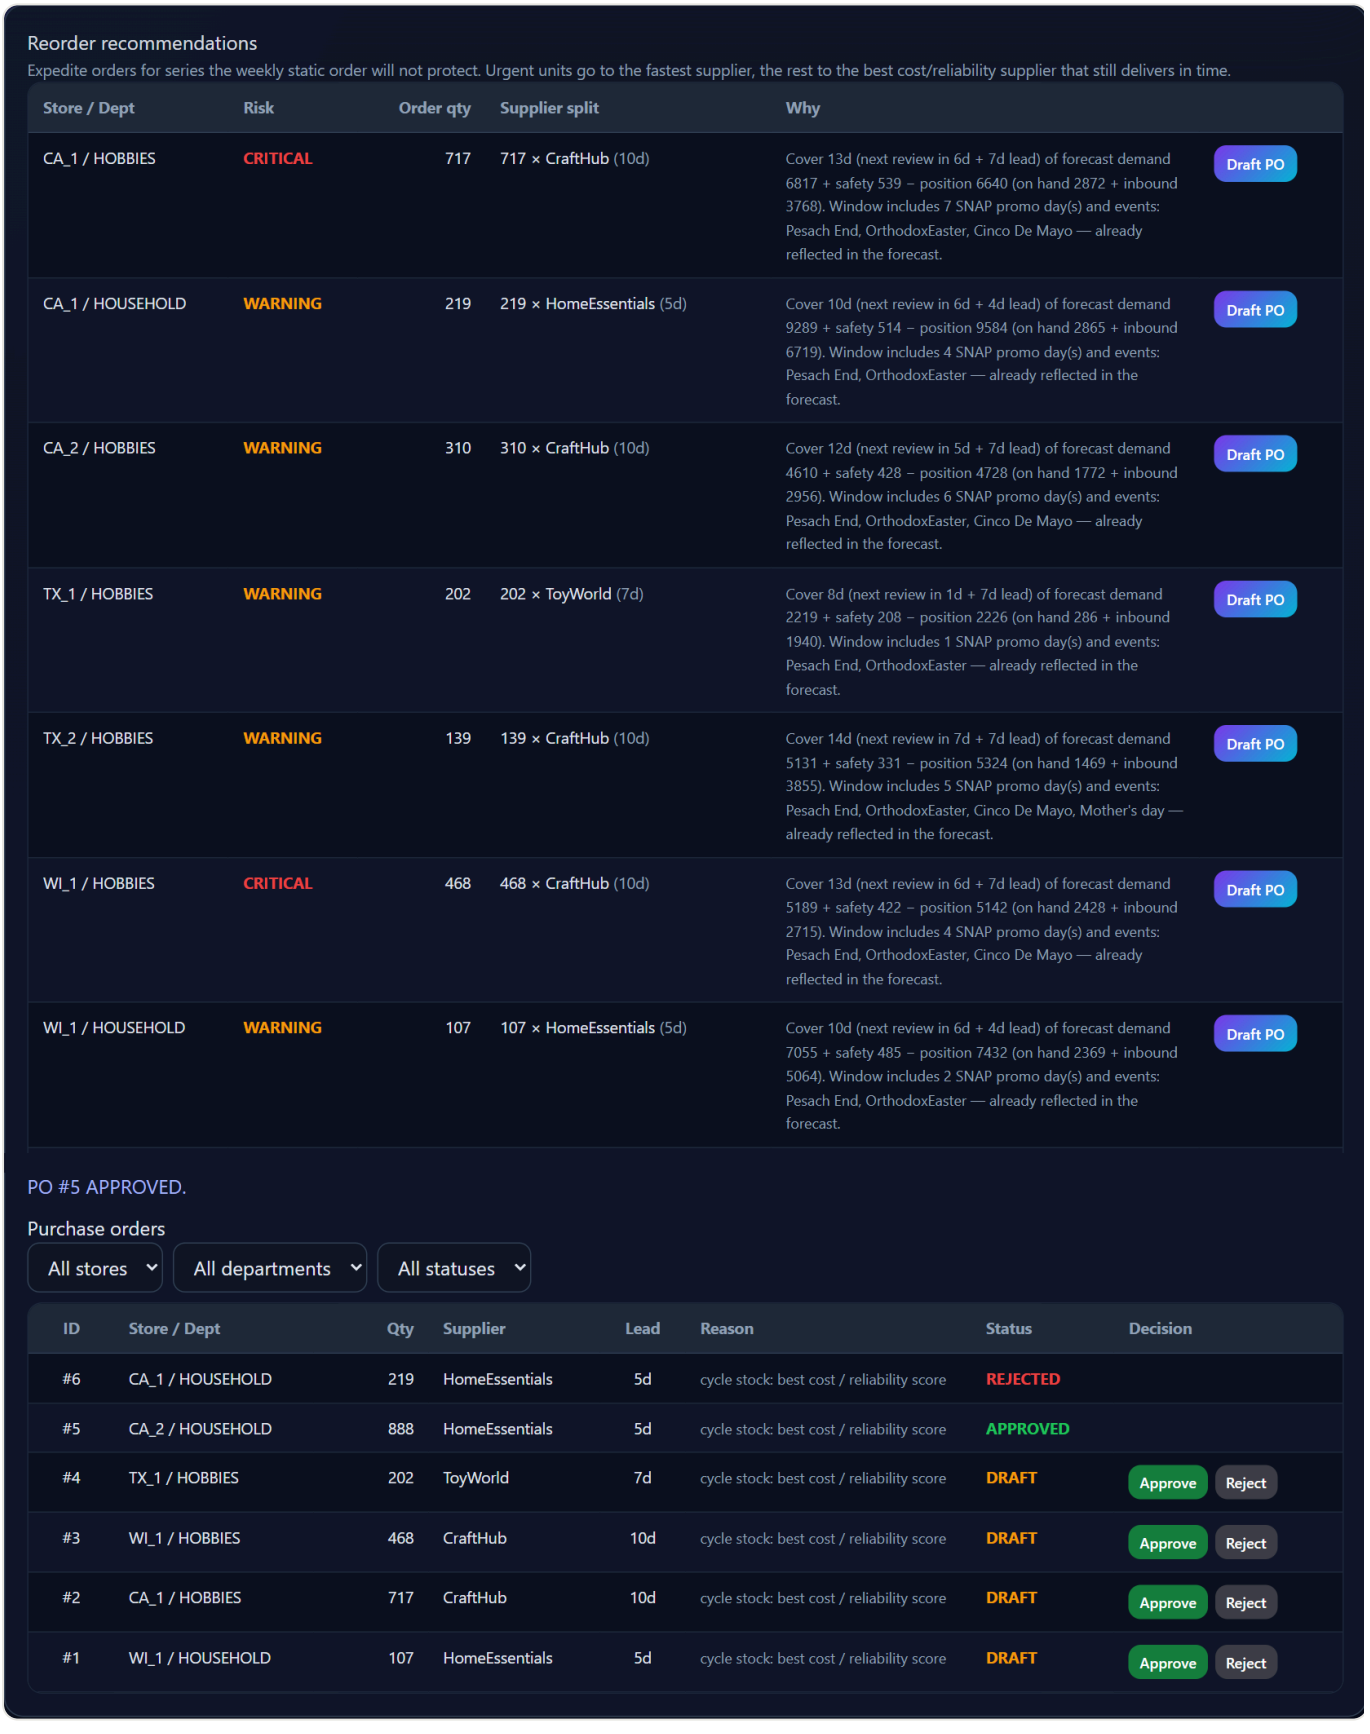
Shaded band: 80% empirical interval from backtest residuals. SNAP promotion days and M5 events are model inputs.



Date	Forecast	Lower bound	Upper bound
2016-04-25	2702.7	2406.5	2906.1
2016-04-26	2431.0	2164.6	2614.0
2016-04-27	2380.0	2119.1	2559.1
2016-04-28	2500.6	2226.5	2688.8
2016-04-29	2979.3	2652.7	3203.5
2016-04-30	3600.3	3205.6	3871.2

Reorders and purchase-order approval

Expedite recommendations with the supplier split and a full rationale, then the purchase orders the agent drafted, which a planner approves or rejects. (The eight OK series, which need no order, are trimmed from this image.)



Reorder recommendations

Expedite orders for series the weekly static order will not protect. Urgent units go to the fastest supplier, the rest to the best cost/reliability supplier that still delivers in time.

Store / Dept	Risk	Order qty	Supplier split	Why	
CA_1 / HOBBIES	CRITICAL	717	717 × CraftHub (10d)	Cover 13d (next review in 6d + 7d lead) of forecast demand 6817 + safety 539 – position 6640 (on hand 2872 + inbound 3768). Window includes 7 SNAP promo day(s) and events: Pesach End, OrthodoxEaster, Cinco De Mayo — already reflected in the forecast.	Draft PO
CA_1 / HOUSEHOLD	WARNING	219	219 × HomeEssentials (5d)	Cover 10d (next review in 6d + 4d lead) of forecast demand 9289 + safety 514 – position 9584 (on hand 2865 + inbound 6719). Window includes 4 SNAP promo day(s) and events: Pesach End, OrthodoxEaster — already reflected in the forecast.	Draft PO
CA_2 / HOBBIES	WARNING	310	310 × CraftHub (10d)	Cover 12d (next review in 5d + 7d lead) of forecast demand 4610 + safety 428 – position 4728 (on hand 1772 + inbound 2956). Window includes 6 SNAP promo day(s) and events: Pesach End, OrthodoxEaster, Cinco De Mayo — already reflected in the forecast.	Draft PO
TX_1 / HOBBIES	WARNING	202	202 × ToyWorld (7d)	Cover 8d (next review in 1d + 7d lead) of forecast demand 2219 + safety 208 – position 2226 (on hand 286 + inbound 1940). Window includes 1 SNAP promo day(s) and events: Pesach End, OrthodoxEaster — already reflected in the forecast.	Draft PO
TX_2 / HOBBIES	WARNING	139	139 × CraftHub (10d)	Cover 14d (next review in 7d + 7d lead) of forecast demand 5131 + safety 331 – position 5324 (on hand 1469 + inbound 3855). Window includes 5 SNAP promo day(s) and events: Pesach End, OrthodoxEaster, Cinco De Mayo, Mother's day — already reflected in the forecast.	Draft PO
WI_1 / HOBBIES	CRITICAL	468	468 × CraftHub (10d)	Cover 13d (next review in 6d + 7d lead) of forecast demand 5189 + safety 422 – position 5142 (on hand 2428 + inbound 2715). Window includes 4 SNAP promo day(s) and events: Pesach End, OrthodoxEaster, Cinco De Mayo — already reflected in the forecast.	Draft PO
WI_1 / HOUSEHOLD	WARNING	107	107 × HomeEssentials (5d)	Cover 10d (next review in 6d + 4d lead) of forecast demand 7055 + safety 485 – position 7432 (on hand 2369 + inbound 5064). Window includes 2 SNAP promo day(s) and events: Pesach End, OrthodoxEaster — already reflected in the forecast.	Draft PO

PO #5 APPROVED.

Purchase orders

All stores | All departments | All statuses

ID	Store / Dept	Qty	Supplier	Lead	Reason	Status	Decision
#6	CA_1 / HOUSEHOLD	219	HomeEssentials	5d	cycle stock: best cost / reliability score	REJECTED	
#5	CA_2 / HOUSEHOLD	888	HomeEssentials	5d	cycle stock: best cost / reliability score	APPROVED	
#4	TX_1 / HOBBIES	202	ToyWorld	7d	cycle stock: best cost / reliability score	DRAFT	Approve Reject
#3	WI_1 / HOBBIES	468	CraftHub	10d	cycle stock: best cost / reliability score	DRAFT	Approve Reject
#2	CA_1 / HOBBIES	717	CraftHub	10d	cycle stock: best cost / reliability score	DRAFT	Approve Reject
#1	WI_1 / HOUSEHOLD	107	HomeEssentials	5d	cycle stock: best cost / reliability score	DRAFT	Approve Reject

9. Demo walkthrough (5 minutes)

1. **Start it:** `python -m uvicorn src.api:app --port 8000`, then `npm --prefix frontend run dev`, and open http://localhost:3000.
2. **Dashboard:** the risk map shows which series run out before their next delivery (2 CRITICAL, 6 WARNING on the M5 data).
3. **Agent:** click "Plan replenishment". Open the reasoning trace to see the agent scan, forecast, check promotions, spikes and Google Trends, then draft supplier-split POs.
4. **Reorders:** approve the drafted POs. Back on the Dashboard, the at-risk count drops, because approved quantities become inbound stock.
5. **Forecast and Report:** change the horizon to see the forecast and its 80% band. The Report tab shows the backtest against both baselines and a downloadable supply-chain report.

Optional: put a free `GROQ_API_KEY` in `.env` and restart the backend. The same questions are then planned by an LLM choosing the tools itself.

10. AI tools leveraged

- **OpenCode (Claude):** scaffolding, refactoring and code review during the build.
- **Claude Code (Claude Opus):** a requirements review and bug-fix pass. It ran the app end to end, then switched to LightGBM, added forecast-driven risk, supplier allocation, the promotion calendar, the replenishment workflow, human PO approval, input validation and test isolation.
- **LightGBM, scikit-learn (metrics), pandas, NumPy, pytrends, FastAPI, Next.js.**
- **LLM planner (optional):** Groq `llama-3.3-70b-versatile` (free tier), OpenAI `gpt-4o-mini`, or a local Ollama model, all through the OpenAI-compatible client.

11. Limitations and next steps

- Inventory and supplier data are simulated or assumed, because M5 has none. A real deployment would read on-hand and inbound stock from the WMS/ERP. The risk engine only needs those inputs.
- The work is at store/department level. Item-level forecasting (M5's 30k series) would need LightGBM with item embeddings or hierarchical reconciliation.
- Google Trends is a weak, US-wide signal. It is reported as context, not used to change order quantities.
- Next steps: price features from `sell_prices.csv`, quantile LightGBM for intervals, a fix for the TX_2/FOODS underperformance, ERP/email export of approved POs, and scheduled daily scans.